

API Reference

Source: docs/api-reference.md

API Reference

Status: internal technical documentation Last updated: 2026-05-23 Extended reference:

API_REFERENCE_EXTENDED.md

FastAPI serves the backend API. User-scoped endpoints require a JWT bearer token in the Authorization: Bearer <token> header unless stated otherwise.

Authentication

POST /auth/login

Authenticates a user with OAuth2PasswordRequestForm.

Request content type: application/x-www-form-urlencoded

Fields:

- username
- password

Response:

```
[json]
{
  "access_token": "string",
  "token_type": "bearer"
}
```

POST /auth/signup

Creates a user account and returns a bearer token. Signup is disabled by default unless SIGNUP_ENABLED=true.

Request body:

```
[json]
{
  "username": "string",
  "email": "optional@example.com",
  "password": "string"
}
```

POST /auth/contact

Accepts a simple contact form payload:

```
[json]
{
  "email": "string",
  "problem": "string"
}
```

Task Status

GET /status/{task_id}

Returns normalized Celery task status. If the task owner was registered in Redis, the endpoint verifies that the authenticated user owns the task.

Response shape:

```
[json]
{
  "task_id": "string",
  "state": "PENDING | STARTED | PROGRESS | RETRY | SUCCESS | FAILURE",
  "progress": 0,
  "stage": "string | null",
  "stage_pl": "string | null",
  "filename": "string | null",
  "attempts": 1,
  "result": {},
  "error": null,
  "updated_at": "ISO-8601"
}
```

User Documents

POST /user/documents/upload

Uploads one user document and starts the parse/chunk/embed Celery pipeline.

Request content type: multipart/form-data

Fields:

- file: PDF, DOCX, XLSX/CSV, TXT or supported text-like file
- tag: optional ESG tag, for example Environmental, Social, Governance

Response:

```
[json]
{
  "task_id": "celery-task-id",
  "status": "queued",
  "message": "Dokument jest przetwarzany w tle. Sprawdź /status/{task_id} po wynik."
}
```

POST /user/documents/delete

Deletes a user document and related chunks.

Request body:

```
[json]
{
  "document_id": "uuid"
}
```

GET /documents/mine

Lists documents for the authenticated user from Supabase.

GET /documents/knowledge

Lists knowledge-base documents. Admin only.

GET /documents/

Lists documents through the document getter router.

Knowledge Base

POST /knowledge/upload

Admin-only multi-file upload to the ESG knowledge base. Each accepted file starts a background parsing task.

Request content type: multipart/form-data

Fields:

- files: one or more files
- tag: optional, defaults to general
- document_type: optional, defaults to general
- version: optional, defaults to 1.0

POST /knowledge/parse-and-store

Admin-only endpoint for parsing a file and storing it in knowledge_documents/knowledge_chunks.

Embeddings

POST /embeddings/generate-for-document

Starts embedding generation for a selected document/chunk table target.

POST /embeddings/generate-for-tag

Starts embedding generation for chunks matching a tag.

POST /embeddings/generate-all

Starts bulk embedding generation.

GET /embeddings/status

Returns basic embedding subsystem status.

Reports

POST /report/generate

Starts asynchronous ESG report generation.

Request body:

```
[json]
{
  "report_scope": "Environmental | Social | Governance | ESG",
  "standard": "GRI | SASB | TCFD"
}
```

standard is optional for backward compatibility. If omitted, the backend uses GRI. The selected standard is injected into the report-generation prompt with the static checklist from backend/report_validation.py.

Response:

```
[json]
{
  "task_id": "celery-task-id",
  "status": "queued",
  "message": "Raport jest generowany w tle. Sprawdz /status/{task_id} po wynik."
}
```

Partial scopes (Environmental, Social, Governance) try tag aliases only. ESG runs without a tag filter. The task result contains:

- status: success or partial_success
- kategoria
- standard: selected reporting standard
- applied_filter
- report_id: stored report id when database persistence succeeds
- used_chunks

- data: structured report JSON or null

The report JSON keeps legacy fields and may include richer fields: `standard_raportowania`, `streszczenie_wykonawcze`, `zakres_i_metodyka`, `szczegolowa_analiza`, `luki_w_danych`, `rekomendacje`, `zgodnosc_ze_standardami`.

GET /report/download/{task_id}

Downloads a generated report as PDF from the cached Celery task result. This does not call the LLM again.

Behavior:

- Requires the task to be in SUCCESS.
- Verifies task ownership when owner metadata exists in Redis.
- Maps task data to ReportData.
- For `partial_success`, renders an empty-state PDF with the message.
- Returns `application/pdf` with Content-Disposition.

POST /report/{report_id}/validate

Validates a stored report against a selected disclosure standard. The validation uses the stored report JSON and parsed `used_chunks`, calls the LLM once, and does not persist the validation result.

Request body:

```
[json]
{
  "standard": "GRI | SASB | TCFD"
}
```

Response:

```
[json]
{
  "status": "success",
  "report_id": "123",
  "standard": "GRI",
  "overall_status": "complete | partial | missing",
  "score": 75,
  "items": [
    {
      "code": "GRI 305-1",
      "label": "Direct Scope 1 GHG emissions",
      "present": true,
      "evidence": "Krotka parafraza z raportu.",
      "recommendation": "Co uzupelnic, jesli brakuje."
    }
  ],
  "summary": "Krotka ocena zgodnosci."
}
```

Supported v1 checklists:

- GRI: GRI 305-1 through 305-5 and GRI 401-1 through 401-3.
- SASB: Engineering & Construction Services metrics, including

IF-EN-160a.1, IF-EN-160a.2, IF-EN-250a.1, IF-EN-250a.2, IF-EN-320a.1, IF-EN-410a.1, IF-EN-410a.2, IF-EN-410b.1, IF-EN-510a.1.

- TCFD: the 11 recommended disclosures across governance, strategy, risk

management, and metrics and targets.

The v1 checklists are static backend definitions, not knowledge-base lookups at request time. The LLM marks individual items as present or missing; the backend then recomputes score and `overall_status` from normalized items so the summary percentage always matches the green/red checklist.

GET /report/{report_id}/validate?standard=GRI

GET alias for the same validation logic. If standard is omitted, GRI is used.

GET /reports/user

Lists report history for the authenticated user.

GET /reports/{report_id}

Returns one stored report owned by the authenticated user, including parsed JSON content and parsed used_chunks.

DELETE /reports/{report_id}

Deletes a stored report owned by the authenticated user.

Chat

POST /chat/ask

Starts a background RAG chat task and stores the user message.

Request body:

```
[json]
{
  "query": "string",
  "tag": "optional tag",
  "session_id": "optional session id"
}
```

Response:

```
[json]
{
  "status": "queued",
  "task_id": "celery-task-id",
  "session_id": "uuid",
  "message": "Pytanie przetwarzane w tle. Sprawdź status zadania."
}
```

GET /chat/sessions

Lists chat sessions for the authenticated user. Supports limit and offset.

GET /chat/sessions/{session_id}/history

Lists messages for a session. Supports limit and offset.

DELETE /chat/sessions/{session_id}

Deletes an owned chat session and its message history.

Utility and Ingestion Endpoints

GET /ping

Health check.

GET /openai-status

Checks whether the backend can initialize OpenAI configuration.

POST /parse

Parses an uploaded file synchronously for parser diagnostics.

POST /process

Legacy async parse-and-store endpoint.

POST /upload

Legacy upload endpoint.

POST /ingest/chunk/url

Fetches a URL, filters/chunks text and returns chunking output.

POST /ingest/chunk/file

Chunks an uploaded file and returns chunking output.

Error Responses

FastAPI validation and error responses use the standard shape:

```
[json]
{
  "detail": "Human-readable error message"
}
```

Common statuses: 400, 401, 403, 404, 409, 413, 422, 500.